

```
1 ##### DESCRIPTION #####
2 # Use to train a model on the specified dataset and hyperparameters.
3
4 import string
5 import random
6 import torch
7 import torch.nn as nn
8 import matplotlib.pyplot as plt
9 import numpy as np
10 import time
11 import sys
12 import os
13 from tqdm import tqdm
14
15 FILE_NAME = 'complete_sherlock_holmes.txt'
16 DATASET = 'Complete Sherlock Holmes'
17
18 ##### HYPERPARAMETERS #####
19 CELL_TYPE = 'RNN' # DEFAULTS TO 'RNN'. Options: ['RNN', 'LSTM', 'GRU', 'RELU']
20 OPTIM_TYPE = 'Adam' # DEFAULTS TO 'Adam'. Options: ['Adam', 'ASGD', 'Adagrad', 'RMSprop']
21 HIDDEN_LAYERS = 1 # DEFAULT: 1
22 HIDDEN_SIZE = 100 # DEFAULT: 100
23
24 LEARNING_RATE = 0.005 # 0.005 for Adam, 0.05 for other optimizers
25 INPUT_SEQUENCE = 128 # DEFAULT: 128
26
27 TRAINING_ITERATIONS = 20000 # DEFAULT: 20000
28 INITIAL_SEQUENCE = '\n' # To use for eval
29
30 print('Running on', CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE)
31
32 all_chars = string.printable
33 n_chars = len(all_chars)
34 file = open('./'+FILE_NAME).read()
35 file_len = len(file)
36
37 def get_random_seq():
38     seq_len = INPUT_SEQUENCE # The length of an input sequence.
39     start_index = random.randint(0, file_len - seq_len)
40     end_index = start_index + seq_len + 1
41     return file[start_index:end_index]
42
43 def seq_to_onehot(seq):
44     tensor = torch.zeros(len(seq), 1, n_chars)
45     for t, char in enumerate(seq):
46         index = all_chars.index(char)
47         tensor[t][0][index] = 1
48     return tensor
49
50 def seq_to_index(seq):
51     tensor = torch.zeros(len(seq), 1)
52     for t, char in enumerate(seq):
53         tensor[t] = all_chars.index(char)
54     return tensor
55
56 def get_input_and_target():
57     seq = get_random_seq()
58     input = seq_to_onehot(seq[:-1]) # Input is represented in one-hot.
59     target = seq_to_index(seq[1:]).long() # Target is represented in index.
```

```

57     return input, target
58
59 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
60 print('Device =', device)
61
62 class Net(nn.Module):
63     def __init__(self):
64         # Initialization.
65         super(Net, self).__init__()
66         self.input_size = n_chars # Input size: Number of unique chars.
67         self.hidden_size = HIDDEN_SIZE
68         self.output_size = n_chars # Output size: Number of unique chars.
69
70         # Ensures the size of the hidden layer stack does not exceed 3
71         self.layers = HIDDEN_LAYERS
72         if HIDDEN_LAYERS>3:
73             self.layers=3
74
75         # Create a rnn cell for the stack
76         def create_cell(size_in, size_out):
77             if CELL_TYPE=='LSTM':
78                 return nn.LSTMCell(size_in, size_out)
79             elif CELL_TYPE=='GRU':
80                 return nn.GRUCell(size_in, size_out)
81             elif CELL_TYPE=='RELU': # Not used in testing report
82                 return nn.RNNCell(size_in, size_out, nonlinearity='relu')
83             else:
84                 return nn.RNNCell(size_in, size_out)
85
86         self.rnn = create_cell(self.input_size, self.hidden_size)
87
88         if HIDDEN_LAYERS>=2:
89             self.rnn2 = create_cell(self.hidden_size, self.hidden_size)
90         if HIDDEN_LAYERS>=3:
91             self.rnn3 = create_cell(self.hidden_size, self.hidden_size)
92
93         self.fc = nn.Linear(self.hidden_size, self.output_size)
94
95     def forward(self, input, hidden, cell, hidden2=False, cell2=False, hidden3=False,
96 cell3=False):
97         # Forward function.
98         # takes in the 'input' and 'hidden' tensors,
99         # can also take in 'cell state' tensor if cell type is 'LSTM',
100        # takes additional hidden and cell state tensors for each layer
101        if CELL_TYPE=='LSTM':
102            hidden, cell = self.rnn(input, (hidden,cell))
103            if self.layers>=2:
104                hidden2, cell2 = self.rnn2(hidden, (hidden2,cell2))
105            if self.layers>=3:
106                hidden3, cell3 = self.rnn3(hidden2, (hidden3,cell3))
107        else:
108            hidden = self.rnn(input, hidden)
109            if self.layers>=2:
110                hidden2 = self.rnn2(hidden, hidden2)
111            if self.layers>=3:
112                hidden3 = self.rnn3(hidden2, hidden3)

```

```

113     # Linear transformation (fully connected layer) to the output
114     if self.layers==3:
115         output = self.fc(hidden3)
116         return output, hidden, cell, hidden2, cell2, hidden3, cell3
117     elif self.layers==2:
118         output = self.fc(hidden2)
119         return output, hidden, cell, hidden2, cell2
120     else:
121         output = self.fc(hidden)
122         return output, hidden, cell
123 def init_hidden(self):
124     # Initial hidden state.
125     return torch.zeros(1, self.hidden_size).to(device)
126 def init_cell(self):
127     # Initial cell state.
128     return torch.zeros(1, self.hidden_size).to(device)
129
130 net = Net()
131 net.to(device)
132
133 # Training step function
134 def train_step(net, opt, input, target):
135     seq_len = input.shape[0]
136     hidden = net.init_hidden() # Initial hidden state
137     cell = net.init_cell() # Initial cell state
138     if HIDDEN_LAYERS >=2:
139         hidden2 = net.init_hidden()
140         cell2 = net.init_cell()
141     if HIDDEN_LAYERS >=3:
142         hidden3 = net.init_hidden()
143         cell3 = net.init_cell()
144
145     net.zero_grad()
146     loss = 0 # Initial loss.
147
148     for t in range(seq_len): # For each one in the input sequence
149         if HIDDEN_LAYERS==3:
150             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(input[t], hidden,
cell, hidden2, cell2, hidden3, cell3)
151         elif HIDDEN_LAYERS==2:
152             output, hidden, cell, hidden2, cell2 = net(input[t], hidden, cell, hidden2,
cell2)
153         else:
154             output, hidden, cell = net(input[t], hidden, cell)
155         loss += loss_func(output, target[t])
156
157     loss.backward() # Backward.
158     opt.step() # Update the weights.
159
160     return loss / seq_len
161
162 # Evaluation step function
163 def eval_step(net, init_seq=INITIAL_SEQUENCE, predicted_len=100):
164     # Initialize the hidden state, input and the predicted sequence
165     hidden = net.init_hidden()
166     cell = net.init_cell()
167     if HIDDEN_LAYERS >=2:

```

```

168         hidden2 = net.init_hidden()
169         cell2 = net.init_cell()
170     if HIDDEN_LAYERS >= 3:
171         hidden3 = net.init_hidden()
172         cell3 = net.init_cell()
173     init_input = seq_to_onehot(init_seq).to(device)
174     predicted_seq = init_seq
175
176     # Use initial string to "build up" hidden state.
177     for t in range(len(init_seq) - 1):
178         if HIDDEN_LAYERS==3:
179             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(init_input[t], hidden,
cell, hidden2, cell2, hidden3, cell3)
180         elif HIDDEN_LAYERS==2:
181             output, hidden, cell, hidden2, cell2 = net(init_input[t], hidden, cell, hidden2,
cell2)
182         else:
183             output, hidden, cell = net(init_input[t], hidden, cell)
184         # Set current input as the last character of the initial string.
185         input = init_input[-1]
186
187     # Predict more characters after the initial string.
188     for t in range(predicted_len):
189         # Get the current output and hidden state.
190         if HIDDEN_LAYERS==3:
191             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(input, hidden, cell,
hidden2, cell2, hidden3, cell3)
192         elif HIDDEN_LAYERS==2:
193             output, hidden, cell, hidden2, cell2 = net(input, hidden, cell, hidden2, cell2)
194         else:
195             output, hidden, cell = net(input, hidden, cell)
196
197     # Sample from the output as a multinomial distribution.
198     try:
199         predicted_index = torch.multinomial(output.view(-1).exp(), 1)[0]
200     except: # Added post to resolve errors with tensors containing 'inf'/'nan' values
201         predicted_index = torch.multinomial(output.view(-1).exp().clamp(0.0, 3.4e38), 1)[0]
202
203     # Add predicted character to the sequence and use it as next input.
204     predicted_char = all_chars[predicted_index]
205     predicted_seq += predicted_char
206     # Use the predicted character to generate the input of next round.
207     input = seq_to_onehot(predicted_char)[0].to(device)
208
209     return predicted_seq
210
211 ##### MAIN ALGORITHM #####
212
213 iters = TRAINING_ITERATIONS # Number of training iterations.
214
215 # The loss variables.
216 all_losses = []
217 # Initialize the optimizer and the loss function.
218 if(OPTIM_TYPE=='ASGD'):
219     opt=torch.optim.ASGD(net.parameters(), lr=LEARNING_RATE)
220 if(OPTIM_TYPE=='Adagrad'): # Not used in testing results

```

```

221     opt=torch.optim.Adagrad(net.parameters(), lr=LEARNING_RATE)
222     if(OPTIM_TYPE=='RMSprop'): # Not used in testing results
223         opt=torch.optim.RMSprop(net.parameters(), lr=LEARNING_RATE)
224     else:
225         opt = torch.optim.Adam(net.parameters(), lr=LEARNING_RATE)
226     loss_func = nn.CrossEntropyLoss()
227
228     # Training procedure.
229     start_time = time.time()
230     for i in tqdm(range(iters)):
231         input, target = get_input_and_target() # Fetch input and target.
232         input, target = input.to(device), target.to(device) # Move to GPU memory.
233         loss = train_step(net, opt, input, target) # Calculate the loss.
234         all_losses.append(loss)
235
236     end_time = time.time()
237     total_time = end_time - start_time
238
239     # Calculates summary of losses
240     i_half = int(len(all_losses)*0.5)
241     i_quart = int(len(all_losses)*0.75)
242     loss_avg = np.sum(np.array(all_losses))/len(all_losses)
243     loss_avg_half = np.sum(np.array(all_losses[i_half:]))/len(all_losses[i_half:])
244     loss_avg_quart = np.sum(np.array(all_losses[i_quart:]))/len(all_losses[i_quart:])
245     loss_list=[elem.item() for elem in [loss_avg,loss_avg_half,loss_avg_quart]]
246     rolling_losses=[]
247     losses_copy = [i.item() for i in all_losses]
248     for i in range(len(losses_copy)):
249         temp=losses_copy[np.max((i-100, 0)):i+1]
250         rolling_losses.append(np.sum(temp)/len(temp))
251
252     plt.xlabel('iters')
253     plt.ylabel('loss')
254     plt.hlines(loss_list,0,len(losses_copy)-1,['red','orange','green'],'dashed')
255     plt.plot(rolling_losses)
256     plt.ylim(0,5)
257
258     print('Avg loss: {}'.format(loss_avg))
259     print('Avg loss last half: {}'.format(loss_avg_half))
260     print('Avg loss last quarter: {}'.format(loss_avg_quart))
261     print()
262     print('Training time: {} sec'.format(total_time))
263     print('{} min | {:.3f} hr'.format(total_time/60, total_time/3600))
264
265     # Creates a folder path to save training results to
266     PATH = 'Results'
267     for folder in [DATASET, CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE]:
268         folder=str(folder)
269         if not os.path.isdir(PATH+'/'+folder):
270             os.mkdir(PATH+'/'+folder)
271         PATH+='/' + folder
272     PATH+='/'
273
274     print()
275     print('Results saved to: {}'.format(PATH))
276     model_path = PATH+'model.pt' # Saves model parameters
277

```

```
278 torch.save(net.state_dict(), model_path)
279
280 # Sequence of all 20000 loss values
281 file = open(PATH+'all_losses.txt', 'w')
282 file.write(' '.join([str(elem) for elem in losses_copy]))
283 file.close()
284
285 # A 5000 char sample generated after training
286 file = open(PATH+'sample.txt', 'w')
287 file.write(eval_step(net, predicted_len=5000))
288 file.close()
289
290 # Information on training
291 file = open(PATH+'info.txt', 'w')
292 file.write('Iterations: {}\n\n'.format(TRAINING_ITERATIONS))
293 file.write('Dataset: {}\nInput Size: {}\nLearning Rate: {}\n\n'.format(DATASET,
INPUT_SEQUENCE, LEARNING_RATE))
294 file.write('Cell Type: {}\nOptimizer: {}\nHidden Layers: {}\nHidden Size:
{}\n\n'.format(CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE))
```

```

1 ##### DESCRIPTION #####
2 # Use to recall a previously trained model to keep generating more text if desired.
3
4 import string
5 import torch
6 import torch.nn as nn
7 import sys
8 import os
9
10 DATASET = 'Complete Sherlock Holmes'
11
12 ##### HYPERPARAMETERS #####
13 # These must be the same as a previously trained model present in the files
14 CELL_TYPE = 'RNN'
15 OPTIM_TYPE = 'Adam'
16 HIDDEN_LAYERS = 1
17 HIDDEN_SIZE = 100
18
19 ##### GENERATION PARAMETERS #####
20 # Alter the length of the new generated sequence and it's initial sequence
21 PREDICTION_LENGTH = 100
22 INITIAL_SEQUENCE = '\n'
23
24 print('Recalling model', CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE)
25 print('Generating {} char, with init_seq: {}'.format(PREDICTION_LENGTH, INITIAL_SEQUENCE))
26
27 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
28 print('Device =', device)
29
30 all_chars = string.printable
31 n_chars = len(all_chars)
32
33 def seq_to_onehot(seq):
34     tensor = torch.zeros(len(seq), 1, n_chars)
35     for t, char in enumerate(seq):
36         index = all_chars.index(char)
37         tensor[t][0][index] = 1
38     return tensor
39
40 class Net(nn.Module):
41     def __init__(self):
42         # Initialization.
43         super(Net, self).__init__()
44         self.input_size = n_chars # Input size: Number of unique chars.
45         self.hidden_size = HIDDEN_SIZE
46         self.output_size = n_chars # Output size: Number of unique chars.
47
48         # Ensures the size of the hidden layer stack does not exceed 3
49         self.layers = HIDDEN_LAYERS
50         if HIDDEN_LAYERS > 3:
51             self.layers = 3
52
53         # Create a rnn cell for the stack
54         def create_cell(size_in, size_out):
55             if CELL_TYPE == 'LSTM':
56                 return nn.LSTMCell(size_in, size_out)

```

```

57         elif CELL_TYPE=='GRU':
58             return nn.GRUCell(size_in, size_out)
59         elif CELL_TYPE=='RELU': # Not used in testing report
60             return nn.RNNCell(size_in, size_out, nonlinearity='relu')
61         else:
62             return nn.RNNCell(size_in, size_out)
63
64     self.rnn = create_cell(self.input_size, self.hidden_size)
65
66     if HIDDEN_LAYERS>=2:
67         self.rnn2 = create_cell(self.hidden_size, self.hidden_size)
68     if HIDDEN_LAYERS>=3:
69         self.rnn3 = create_cell(self.hidden_size, self.hidden_size)
70
71     self.linear = nn.Linear(self.hidden_size, self.output_size)
72
73     def forward(self, input, hidden, cell, hidden2=False, cell2=False, hidden3=False,
74 cell3=False):
75         # Forward function.
76         # takes in the 'input' and 'hidden' tensors,
77         # can also take in 'cell state' tensor if cell type is 'LSTM',
78         # takes additional hidden and cell state tensors for each layer
79         if CELL_TYPE=='LSTM':
80             hidden, cell = self.rnn(input, (hidden,cell))
81             if self.layers>=2:
82                 hidden2, cell2 = self.rnn2(hidden, (hidden2,cell2))
83             if self.layers>=3:
84                 hidden3, cell3 = self.rnn3(hidden2, (hidden3,cell3))
85         else:
86             hidden = self.rnn(input, hidden)
87             if self.layers>=2:
88                 hidden2 = self.rnn2(hidden, hidden2)
89             if self.layers>=3:
90                 hidden3 = self.rnn3(hidden2, hidden3)
91
92         # Linear transformation (fully connected layer) to the output
93         if self.layers==3:
94             output = self.linear(hidden3)
95             return output, hidden, cell, hidden2, cell2, hidden3, cell3
96         elif self.layers==2:
97             output = self.linear(hidden2)
98             return output, hidden, cell, hidden2, cell2
99         else:
100             output = self.linear(hidden)
101             return output, hidden, cell
102
103     def init_hidden(self):
104         # Initial hidden state.
105         return torch.zeros(1, self.hidden_size).to(device)
106
107     def init_cell(self):
108         # Initial cell state.
109         return torch.zeros(1, self.hidden_size).to(device)
110
111     def eval_step(net, init_seq='\n', predicted_len=100):
112         # Initialize the hidden state, input and the predicted sequence
113         hidden = net.init_hidden()
114         cell = net.init_cell()
115         if HIDDEN_LAYERS >=2:

```



```

113         hidden2 = net.init_hidden()
114         cell2 = net.init_cell()
115     if HIDDEN_LAYERS >= 3:
116         hidden3 = net.init_hidden()
117         cell3 = net.init_cell()
118     init_input = seq_to_onehot(init_seq).to(device)
119     predicted_seq = init_seq
120
121     # Use initial string to "build up" hidden state.
122     for t in range(len(init_seq) - 1):
123         if HIDDEN_LAYERS==3:
124             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(init_input[t], hidde
cell, hidden2, cell2, hidden3, cell3)
125         elif HIDDEN_LAYERS==2:
126             output, hidden, cell, hidden2, cell2 = net(init_input[t], hidden, cell, hidden2
cell2)
127         else:
128             output, hidden, cell = net(init_input[t], hidden, cell)
129     # Set current input as the last character of the initial string.
130     input = init_input[-1]
131
132     # Predict more characters after the initial string.
133     for t in range(predicted_len):
134         # Get the current output and hidden state.
135         if HIDDEN_LAYERS==3:
136             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(input, hidden, cell,
hidden2, cell2, hidden3, cell3)
137         elif HIDDEN_LAYERS==2:
138             output, hidden, cell, hidden2, cell2 = net(input, hidden, cell, hidden2, cell2)
139         else:
140             output, hidden, cell = net(input, hidden, cell)
141
142         # Sample from the output as a multinomial distribution.
143         try:
144             predicted_index = torch.multinomial(output.view(-1).exp(), 1)[0]
145         except: # Added post to resolve errors with tensors containing 'inf'/'nan' values
146             predicted_index = torch.multinomial(output.view(-1).exp().clamp(0.0, 3.4e38), 1)
[0]
147         # Add predicted character to the sequence and use it as next input.
148         predicted_char = all_chars[predicted_index]
149         predicted_seq += predicted_char
150         # Use the predicted character to generate the input of next round.
151         input = seq_to_onehot(predicted_char)[0].to(device)
152
153     return predicted_seq
154
155 PATH = 'Results'
156 for folder in [DATASET, CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE]:
157     folder=str(folder)
158     PATH+='/' + folder
159 PATH+=' /model.pt'
160
161 net=Net()
162 net.to(device)
163
164 net.load_state_dict(torch.load(PATH))
165 net.eval()

```

```
166  
167 print(eval_step(net, init_seq=INITIAL_SEQUENCE, predicted_len=PREDICTION_LENGTH))
```

Analysis of Results

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import re
import string
```

Calculations about expected metrics

Here I have taken the dataset and determined what metrics may be expected out of some randomly picked 5000 character sequence to compare with the generated sample from each model.

```
In [2]: file = open('./complete_sherlock_holmes.txt').read()
file_len = len(file)
```

```
In [3]: words = re.split(' |\n',file)
word_list = [w for w in words if len(w)>0]
n1 = 5000*len(word_list)/file_len
print('Expected words per 5000 char: {}'.format(n1))
```

Expected words per 5000 char: 849.9241978679597

This would be the expected "word count" out of some 5000 character sequence. That is, the number of separate sequences of alphanumeric characters and symbols, upper or lowercase, separated by spaces or new lines.

```
In [4]: #string.printable[:95]
file2 = ''.join([i for i in file if i in string.printable[:95]])
n2 = len(file2)/file_len*5000
print('Expected characters w/ spaces per 5000 char: {}'.format(n2))
file3 = ''.join([i for i in file if i in string.printable[:94]])
n3 = len(file3)/file_len*5000
print('Expected characters w/o spaces per 5000 char: {}'.format(n3))
```

Expected characters w/ spaces per 5000 char: 4901.287202110336

Expected characters w/o spaces per 5000 char: 3732.427788378544


Characters with spaces would be the 5000 characters minus any format characters such as new line ('\n'), and without spaces would also subtract the number of space characters.

Tracking of Losses

For each model I tracked the losses across all iterations in order to make sure that the model could converge after 20000 iterations. In order to smooth out the graph and large jumps in the loss, I took the rolling loss from the previous 100 iterations and plotted this sequence of losses along with horizontal lines showing the average loss across all iterations, across the last half (10000 iterations), and last quarter (5000 iterations). Shown in dotted red, orange, and green lines respectively.

```
In [5]: file = open('./Results/Complete Sherlock Holmes/RNN/Adam/1/100/all_losses.txt').r
losses = [float(i) for i in file.split(' ')]
rolling_losses=[]
for i in range(len(losses)):
    temp=losses[np.max((i-100, 0)):i+1]
    rolling_losses.append(np.sum(temp)/len(temp))
avg = np.sum(losses)/len(losses)
avg_half = np.sum(losses[10000:])/len(losses[10000:])
avg_quart = np.sum(losses[5000:])/len(losses[5000:])
plt.plot(rolling_losses)
plt.hlines([avg,avg_half,avg_quart],0,len(losses)-1,['red','orange','green'],'dashdot')
plt.title('Rolling losses for: RNN, Adam, 1, 100')
plt.xlabel('Iterations')

plt.savefig('Images/loss_graph.png')
plt.show()
```



(Above) Example plot showing the rolling losses while training the model with cell type: RNN, optimizer: Adam, 1 hidden layer and layer size of 100.

```

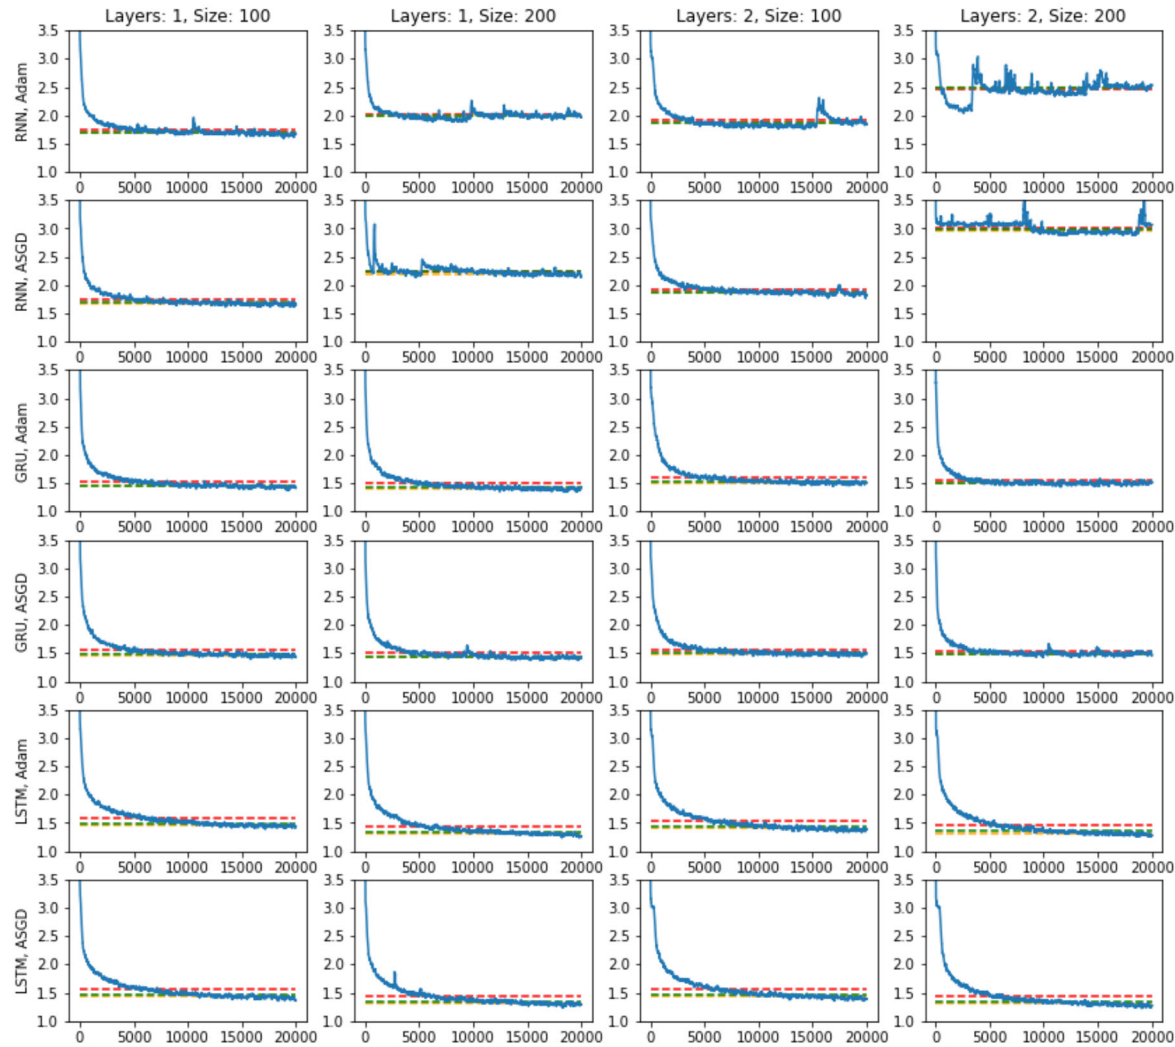
In [6]: plt.figure(figsize=(14,13))
sub_count = 1
for cell_type in ['RNN', 'GRU', 'LSTM']:
    for optim in ['Adam', 'ASGD']:
        for numlayers in ['1', '2']:
            for size in ['100', '200']:
                file_name = './Results/Complete Sherlock Holmes/'
                file_name += '{}/{}/{}/all_losses.txt'.format(cell_type, optim, size)
                file = open(file_name).read()
                losses = [float(i) for i in file.split(' ')]
                rolling_losses=[]
                for i in range(len(losses)):
                    temp=losses[np.max((i-100, 0)):i+1]
                    rolling_losses.append(np.sum(temp)/len(temp))

                avg = np.sum(losses)/len(losses)
                avg_half = np.sum(losses[10000:])/len(losses[10000:])
                avg_quart = np.sum(losses[5000:])/len(losses[5000:])

                plt.subplot(6,4,sub_count)
                plt.plot(rolling_losses)
                plt.hlines([avg,avg_half,avg_quart],0,len(losses)-1,['red','orange','blue'])
                plt.ylim(1.0,3.5)
                if sub_count <= 4:
                    plt.title('Layers: {}, Size: {}'.format(numlayers, size))
                if sub_count%4 == 1:
                    plt.ylabel('{} {}'.format(cell_type, optim))

                sub_count+=1
plt.savefig('Images/all_loss_graphs.png')

```



(Above) Rolling losses for all models


Table of performance results for each model

```
In [7]: df = pd.read_csv('results.csv', sep=',')
df
```

Out[7]:

	Cell Type	Optimizer	# Hidden Layers	Hidden Layer Size	Avg Loss	last half	last quarter	training time (min)	Sample char count w/o spaces	Sample char count w/ spaces
0	RNN	Adam	1	100	1.753240	1.690466	1.673080	29.160343	3744	4892
1	RNN	Adam	1	200	2.011810	2.003207	1.994535	29.708095	3864	4892
2	RNN	Adam	2	100	1.908531	1.872844	1.929789	43.287523	3769	4881
3	RNN	Adam	2	200	2.457533	2.482899	2.530376	48.069504	4062	4912
4	RNN	ASGD	1	100	1.745751	1.673216	1.662270	43.538487	3785	4916
5	RNN	ASGD	1	200	2.252878	2.204309	2.192678	29.156702	3690	4867
6	RNN	ASGD	2	100	1.932660	1.863559	1.854473	67.390463	3546	4848
7	RNN	ASGD	2	200	3.019652	2.958157	2.986661	45.852272	3898	4960
8	GRU	Adam	1	100	1.530559	1.448299	1.436921	27.283551	3713	4910
9	GRU	Adam	1	200	1.494480	1.410895	1.401998	41.744309	3726	4918
10	GRU	Adam	2	100	1.609544	1.515976	1.507607	44.035200	3649	4868
11	GRU	Adam	2	200	1.546415	1.498690	1.502236	63.860197	3710	4855
12	GRU	ASGD	1	100	1.552805	1.468210	1.456479	41.326073	3715	4890
13	GRU	ASGD	1	200	1.505433	1.428692	1.419809	29.475967	3080	4642
14	GRU	ASGD	2	100	1.568220	1.488219	1.484442	66.975609	3653	4884
15	GRU	ASGD	2	200	1.537782	1.493338	1.489142	43.401833	3631	4893
16	LSTM	Adam	1	100	1.581827	1.464306	1.449594	26.358542	3807	4918
17	LSTM	Adam	1	200	1.435656	1.317489	1.297824	28.450239	3708	4912
18	LSTM	Adam	2	100	1.533304	1.403713	1.384826	59.830156	3715	4907
19	LSTM	Adam	2	200	1.458838	1.320205	1.305445	43.398384	3698	4895
20	LSTM	ASGD	1	100	1.554531	1.434153	1.417843	28.503249	3555	4854
21	LSTM	ASGD	1	200	1.443373	1.324596	1.304677	27.946114	3767	4910
22	LSTM	ASGD	2	100	1.575714	1.433288	1.413820	58.016687	3772	4912
23	LSTM	ASGD	2	200	1.447734	1.305508	1.286445	60.512118	3764	4913

```
In [8]: plt.figure(figsize=(12,15))
sub_count = 1
for cell_type in ['RNN', 'GRU', 'LSTM']:
    for optim in ['Adam', 'ASGD']:
        df_sub = df
        df_sub = df_sub[df_sub['Cell Type']==cell_type]
        df_sub = df_sub[df_sub['Optimizer']==optim]
        df_sub = df_sub[['Avg Loss', 'last half', 'last quarter']]
        losses = df_sub.to_numpy()
        plt.subplot(4,2,sub_count)
        avgs = losses[:,0]
        halves = losses[:,1]
        quarts = losses[:,2]
        ind = np.arange(4)
        width = 0.25
        plt.bar(ind, avgs, width, label='Avg loss (all)')
        plt.bar(ind + width, halves, width, label='(last half)')
        plt.bar(ind + 2*width, halves, width, label='(last quarter)')
        plt.ylim(1.0,3.1)
        if sub_count%2 == 1:
            plt.ylabel('Cell Type: {}'.format(cell_type))
        plt.xlabel('(Hidden Layers, Hidden Layer Size)')
        if sub_count <=2 :
            plt.title('Optimizer: {}'.format(optim))
        plt.xticks(ind + width, ('1, 100', '1, 200', '2, 100', '2, 200'))
        plt.legend(loc='best')
        sub_count+=1
plt.savefig('Images/loss_comparison.png')
```





(Above) Bar chart comparison of average loss per all losses, last half losses, and last quarter losses.

```
In [9]: labels = []
for cell_type in ['RNN', 'GRU', 'LSTM']:
    for optim in ['Adam', 'ASGD']:
        for numlayers in ['1', '2']:
            for size in ['100', '200']:
                labels.append('{} {}, {}, {}'.format(cell_type, optim, numlayers, size))
```

```
In [10]: losses = df['last quarter'].to_numpy()
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(losses)),losses,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.ylim(1.0,3.1)
plt.vlines([7.5,15.5], 1.0, 3.05, 'gray')
plt.vlines([3.5,11.5,19.5], 1.0, 3.05, 'gray','dashed')
plt.title('Mean cross entropy loss of last 5,000 iterations')
plt.tight_layout()
plt.savefig('Images/loss_last_quarter.png')
plt.show()

print('Lowest cross entropy loss: LSTM, ASGD, 2, 200. ({}).format(df['last quarter'])
```




Lowest cross entropy loss: LSTM, ASGD, 2, 200. (1.286444902)

(Above) Comparison of last quarter cross entropy loss across all models.

```
In [11]: times = df['training time (min)'].to_numpy()
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(times)),times,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], 0, 70, 'gray')
plt.vlines([3.5,11.5,19.5], 0, 70, 'gray','dashed')
plt.ylabel('Training time (min)')
plt.title('Training time for 20,000 iterations')
plt.tight_layout()
plt.savefig('Images/training_time.png')
plt.show()


print('Shortest training time: LSTM, Adam, 1, 100. ({} min)'.format(df['training
```



Shortest training time: LSTM, Adam, 1, 100. (26.35854196 min)

```
In [12]: errors = df['% spelling errors'].to_numpy()*100
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(times)),errors,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], 0, 75, 'gray')
plt.vlines([3.5,11.5,19.5], 0, 75, 'gray','dashed')
plt.ylabel('# Spelling errors / word count (%)')
plt.title('Percentage of incorrectly spelled words in 5,000 char sequence')
plt.tight_layout()
plt.savefig('Images/spelling_percentage.png')
plt.show()


print('Smallest percentage of incorrectly spelled words: LSTM, ASGD, 2, 200. ({}%
```



Smallest percentage of incorrectly spelled words: LSTM, ASGD, 2, 200. (14.7161066%)


```
In [13]: gerrors = df['grammer errors'].to_numpy()
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(times)),gerrors,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], 0, 60, 'gray')
plt.vlines([3.5,11.5,19.5], 0, 60, 'gray','dashed')
plt.title('Grammatical errors in 5,000 char sample sequence')
plt.tight_layout()
plt.savefig('Images/grammar_errors.png')
plt.show()

print('Fewest grammatical errors: GRU, Adam, 2, 200. ({} )'.format(df['grammer err
```



Fewest grammatical errors: GRU, Adam, 2, 200. (11)

```
In [14]: data = (df['Sample word count'].to_numpy() - n1)**-1
plt.figure(figsize=(10,7))
colors=['g' if i>=0 else 'r' for i in data]
plt.bar(np.arange(len(times)),data,color=colors,tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], data.min(), data.max(), 'gray')
plt.vlines([3.5,11.5,19.5], data.min(), data.max(), 'gray','dashed')
plt.ylabel('1 / (Sample word count - expected word count)')
plt.title('Inverse difference between expected and actual word count (per 5,000 c
plt.tight_layout()
plt.savefig('Images/word_count_comparison.png')
plt.show()
```




```
In [15]: plt.figure(figsize=(16,6))
plt.subplot(1,2,1)
data = (df['Sample char count w/o spaces'].to_numpy() - n3)**-1

colors=['g' if i>=0 else 'r' for i in data]
plt.bar(np.arange(len(times)),data,color=colors,tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], -0.16, 0.18, 'gray')
plt.vlines([3.5,11.5,19.5], -0.16, 0.18, 'gray','dashed')
plt.ylabel('1 / (Sample char count - expected char count)')
plt.title('Inverse difference between expected and actual\ncharacter count w/o spaces')

plt.subplot(1,2,2)
data = (df['Sample char count w/ spaces'].to_numpy() - n2)**-1

colors=['g' if i>=0 else 'r' for i in data]
plt.bar(np.arange(len(times)),data,color=colors,tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], -0.16, 0.18, 'gray')
plt.vlines([3.5,11.5,19.5], -0.16, 0.18, 'gray','dashed')
plt.title('Inverse difference between expected and actual\ncharacter count w/ spaces')

plt.savefig('Images/char_count_comparison.png')
plt.show()
```



In []:

