

# Analysis of Results

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import re
import string
```

## Calculations about expected metrics

Here I have taken the dataset and determined what metrics may be expected out of some randomly picked 5000 character sequence to compare with the generated sample from each model.

```
In [2]: file = open('./complete_sherlock_holmes.txt').read()
file_len = len(file)
```

```
In [3]: words = re.split(' |\n',file)
word_list = [w for w in words if len(w)>0]
n1 = 5000*len(word_list)/file_len
print('Expected words per 5000 char: {}'.format(n1))
```

Expected words per 5000 char: 849.9241978679597

This would be the expected "word count" out of some 5000 character sequence. That is, the number of separate sequences of alphanumeric characters and symbols, upper or lowercase, separated by spaces or new lines.

```
In [4]: #string.printable[:95]
file2 = ''.join([i for i in file if i in string.printable[:95]])
n2 = len(file2)/file_len*5000
print('Expected characters w/ spaces per 5000 char: {}'.format(n2))
file3 = ''.join([i for i in file if i in string.printable[:94]])
n3 = len(file3)/file_len*5000
print('Expected characters w/o spaces per 5000 char: {}'.format(n3))
```

Expected characters w/ spaces per 5000 char: 4901.287202110336

Expected characters w/o spaces per 5000 char: 3732.427788378544

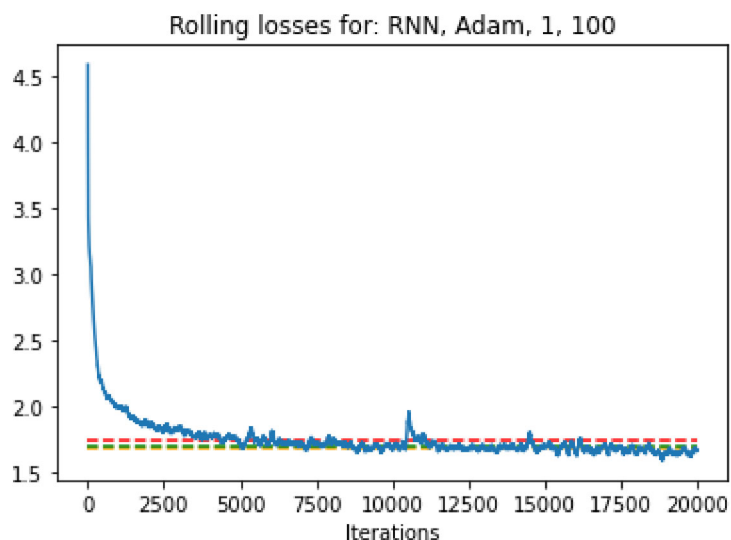
Characters with spaces would be the 5000 characters minus any format characters such as new line ('\n'), and without spaces would also subtract the number of space characters.

## Tracking of Losses

For each model I tracked the losses across all iterations in order to make sure that the model could converge after 20000 iterations. In order to smooth out the graph and large jumps in the loss, I took the rolling loss from the previous 100 iterations and plotted this sequence of losses along with horizontal lines showing the average loss across all iterations, across the last half (10000 iterations), and last quarter (5000 iterations). Shown in dotted red, orange, and green lines respectively.

```
In [5]: file = open('./Results/Complete Sherlock Holmes/RNN/Adam/1/100/all_losses.txt').r
losses = [float(i) for i in file.split(' ')]
rolling_losses=[]
for i in range(len(losses)):
    temp=losses[np.max((i-100, 0)):i+1]
    rolling_losses.append(np.sum(temp)/len(temp))
avg = np.sum(losses)/len(losses)
avg_half = np.sum(losses[10000:])/len(losses[10000:])
avg_quart = np.sum(losses[5000:])/len(losses[5000:])
plt.plot(rolling_losses)
plt.hlines([avg,avg_half,avg_quart],0,len(losses)-1,['red','orange','green'],'das
plt.title('Rolling losses for: RNN, Adam, 1, 100')
plt.xlabel('Iterations')

plt.savefig('Images/loss_graph.png')
plt.show()
```



(Above) Example plot showing the rolling losses while training the model with cell type: RNN, optimizer: Adam, 1 hidden layer and layer size of 100.

```

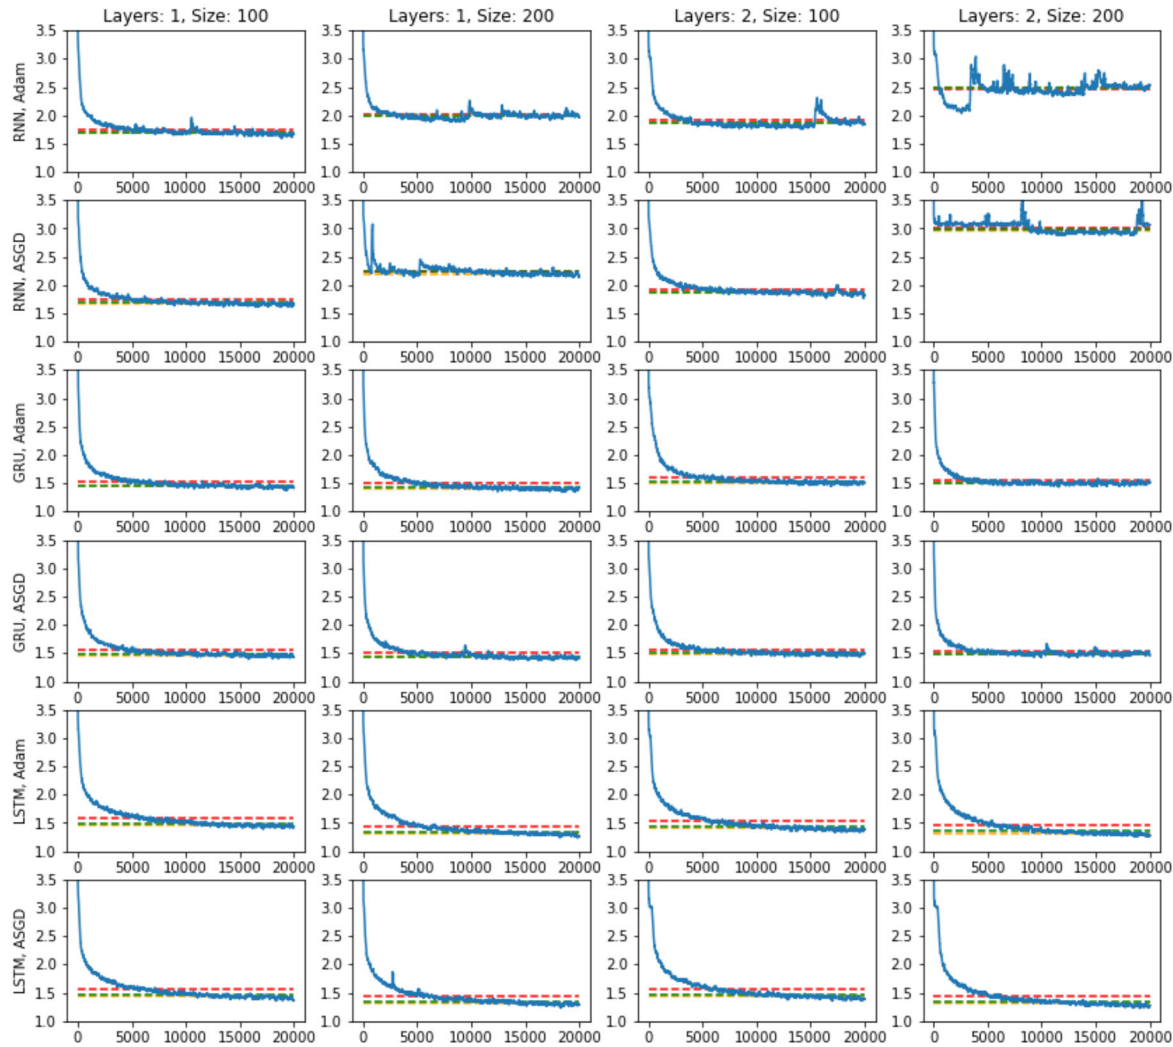
In [6]: plt.figure(figsize=(14,13))
sub_count = 1
for cell_type in ['RNN', 'GRU', 'LSTM']:
    for optim in ['Adam', 'ASGD']:
        for numlayers in ['1', '2']:
            for size in ['100', '200']:
                file_name = './Results/Complete Sherlock Holmes/'
                file_name += '{}/{}/{}/all_losses.txt'.format(cell_type, optim, size)
                file = open(file_name).read()
                losses = [float(i) for i in file.split(' ')]
                rolling_losses=[]
                for i in range(len(losses)):
                    temp=losses[np.max((i-100, 0)):i+1]
                    rolling_losses.append(np.sum(temp)/len(temp))

                avg = np.sum(losses)/len(losses)
                avg_half = np.sum(losses[10000:])/len(losses[10000:])
                avg_quart = np.sum(losses[5000:])/len(losses[5000:])

                plt.subplot(6,4,sub_count)
                plt.plot(rolling_losses)
                plt.hlines([avg,avg_half,avg_quart],0,len(losses)-1,['red','orange','green'])
                plt.ylim(1.0,3.5)
                if sub_count <= 4:
                    plt.title('Layers: {}, Size: {}'.format(numlayers, size))
                if sub_count%4 == 1:
                    plt.ylabel('{} , {}'.format(cell_type, optim))

                sub_count+=1
plt.savefig('Images/all_loss_graphs.png')

```



(Above) Rolling losses for all models


Table of performance results for each model

```
In [7]: df = pd.read_csv('results.csv', sep=',')
df
```

Out[7]:

	Cell Type	Optimizer	# Hidden Layers	Hidden Layer Size	Avg Loss	last half	last quarter	training time (min)	Sample char count w/o spaces	Sample char count w/ spaces
0	RNN	Adam	1	100	1.753240	1.690466	1.673080	29.160343	3744	4892
1	RNN	Adam	1	200	2.011810	2.003207	1.994535	29.708095	3864	4892
2	RNN	Adam	2	100	1.908531	1.872844	1.929789	43.287523	3769	4881
3	RNN	Adam	2	200	2.457533	2.482899	2.530376	48.069504	4062	4912
4	RNN	ASGD	1	100	1.745751	1.673216	1.662270	43.538487	3785	4916
5	RNN	ASGD	1	200	2.252878	2.204309	2.192678	29.156702	3690	4867
6	RNN	ASGD	2	100	1.932660	1.863559	1.854473	67.390463	3546	4848
7	RNN	ASGD	2	200	3.019652	2.958157	2.986661	45.852272	3898	4960
8	GRU	Adam	1	100	1.530559	1.448299	1.436921	27.283551	3713	4910
9	GRU	Adam	1	200	1.494480	1.410895	1.401998	41.744309	3726	4918
10	GRU	Adam	2	100	1.609544	1.515976	1.507607	44.035200	3649	4868
11	GRU	Adam	2	200	1.546415	1.498690	1.502236	63.860197	3710	4855
12	GRU	ASGD	1	100	1.552805	1.468210	1.456479	41.326073	3715	4890
13	GRU	ASGD	1	200	1.505433	1.428692	1.419809	29.475967	3080	4642
14	GRU	ASGD	2	100	1.568220	1.488219	1.484442	66.975609	3653	4884
15	GRU	ASGD	2	200	1.537782	1.493338	1.489142	43.401833	3631	4893
16	LSTM	Adam	1	100	1.581827	1.464306	1.449594	26.358542	3807	4918
17	LSTM	Adam	1	200	1.435656	1.317489	1.297824	28.450239	3708	4912
18	LSTM	Adam	2	100	1.533304	1.403713	1.384826	59.830156	3715	4907
19	LSTM	Adam	2	200	1.458838	1.320205	1.305445	43.398384	3698	4895
20	LSTM	ASGD	1	100	1.554531	1.434153	1.417843	28.503249	3555	4854
21	LSTM	ASGD	1	200	1.443373	1.324596	1.304677	27.946114	3767	4910
22	LSTM	ASGD	2	100	1.575714	1.433288	1.413820	58.016687	3772	4912
23	LSTM	ASGD	2	200	1.447734	1.305508	1.286445	60.512118	3764	4913

```
In [8]: plt.figure(figsize=(12,15))
sub_count = 1
for cell_type in ['RNN', 'GRU', 'LSTM']:
    for optim in ['Adam', 'ASGD']:
        df_sub = df
        df_sub = df_sub[df_sub['Cell Type']==cell_type]
        df_sub = df_sub[df_sub['Optimizer']==optim]
        df_sub = df_sub[['Avg Loss', 'last half', 'last quarter']]
        losses = df_sub.to_numpy()
        plt.subplot(4,2,sub_count)
        avgs = losses[:,0]
        halves = losses[:,1]
        quarts = losses[:,2]
        ind = np.arange(4)
        width = 0.25
        plt.bar(ind, avgs, width, label='Avg loss (all)')
        plt.bar(ind + width, halves, width, label='(last half)')
        plt.bar(ind + 2*width, halves, width, label='(last quarter)')
        plt.ylim(1.0,3.1)
        if sub_count%2 == 1:
            plt.ylabel('Cell Type: {}'.format(cell_type))
        plt.xlabel('(Hidden Layers, Hidden Layer Size)')
        if sub_count <=2 :
            plt.title('Optimizer: {}'.format(optim))
        plt.xticks(ind + width, ('1, 100', '1, 200', '2, 100', '2, 200'))
        plt.legend(loc='best')
        sub_count+=1
plt.savefig('Images/loss_comparison.png')
```

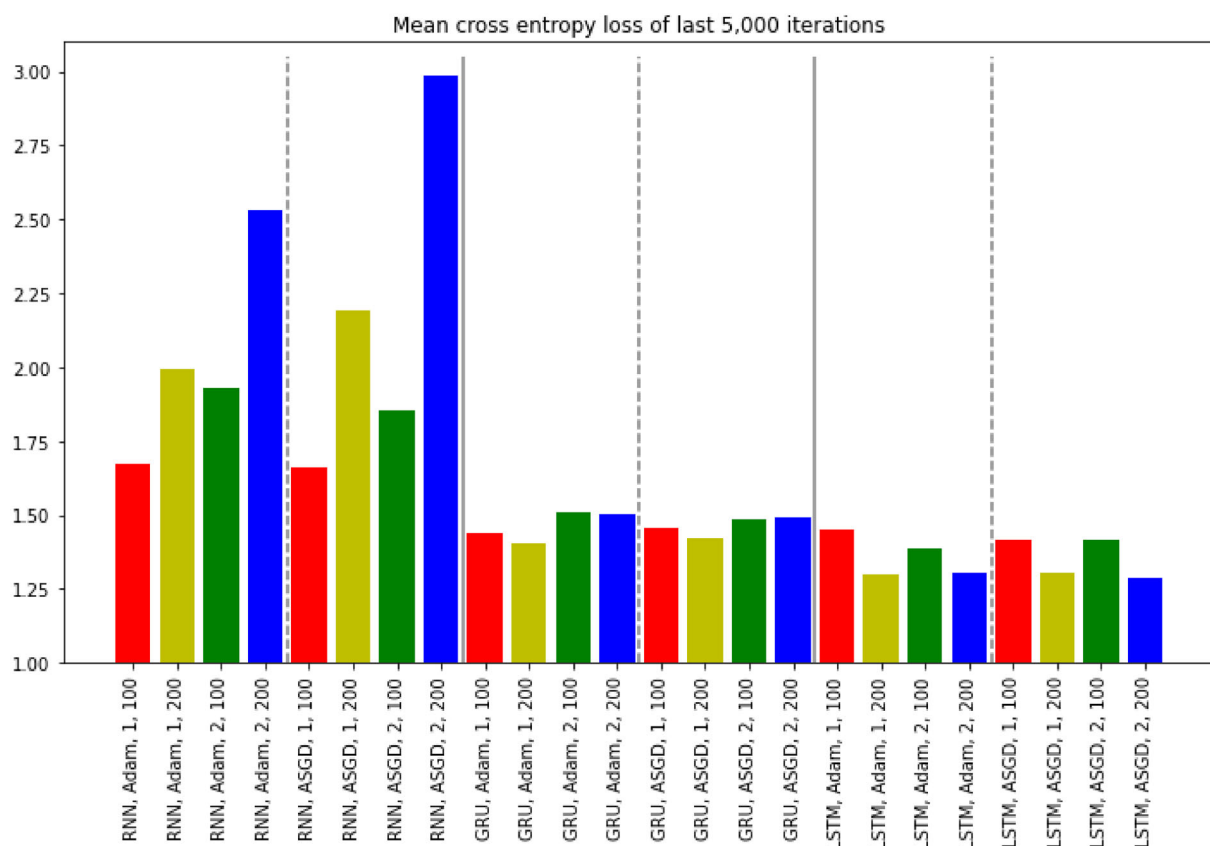


(Above) Bar chart comparison of average loss per all losses, last half losses, and last quarter losses.

```
In [9]: labels = []
for cell_type in ['RNN', 'GRU', 'LSTM']:
    for optim in ['Adam', 'ASGD']:
        for numlayers in ['1', '2']:
            for size in ['100', '200']:
                labels.append('{} {}, {}, {}'.format(cell_type, optim, numlayers, size))
```

```
In [10]: losses = df['last quarter'].to_numpy()
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(losses)),losses,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.ylim(1.0,3.1)
plt.vlines([7.5,15.5], 1.0, 3.05, 'gray')
plt.vlines([3.5,11.5,19.5], 1.0, 3.05, 'gray','dashed')
plt.title('Mean cross entropy loss of last 5,000 iterations')
plt.tight_layout()
plt.savefig('Images/loss_last_quarter.png')
plt.show()

print('Lowest cross entropy loss: LSTM, ASGD, 2, 200. ({}).format(df['last quarter'])
```




Lowest cross entropy loss: LSTM, ASGD, 2, 200. (1.286444902)



(Above) Comparison of last quarter cross entropy loss across all models.

```
In [11]: times = df['training time (min)'].to_numpy()
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(times)),times,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], 0, 70, 'gray')
plt.vlines([3.5,11.5,19.5], 0, 70, 'gray','dashed')
plt.ylabel('Training time (min)')
plt.title('Training time for 20,000 iterations')
plt.tight_layout()
plt.savefig('Images/training_time.png')
plt.show()

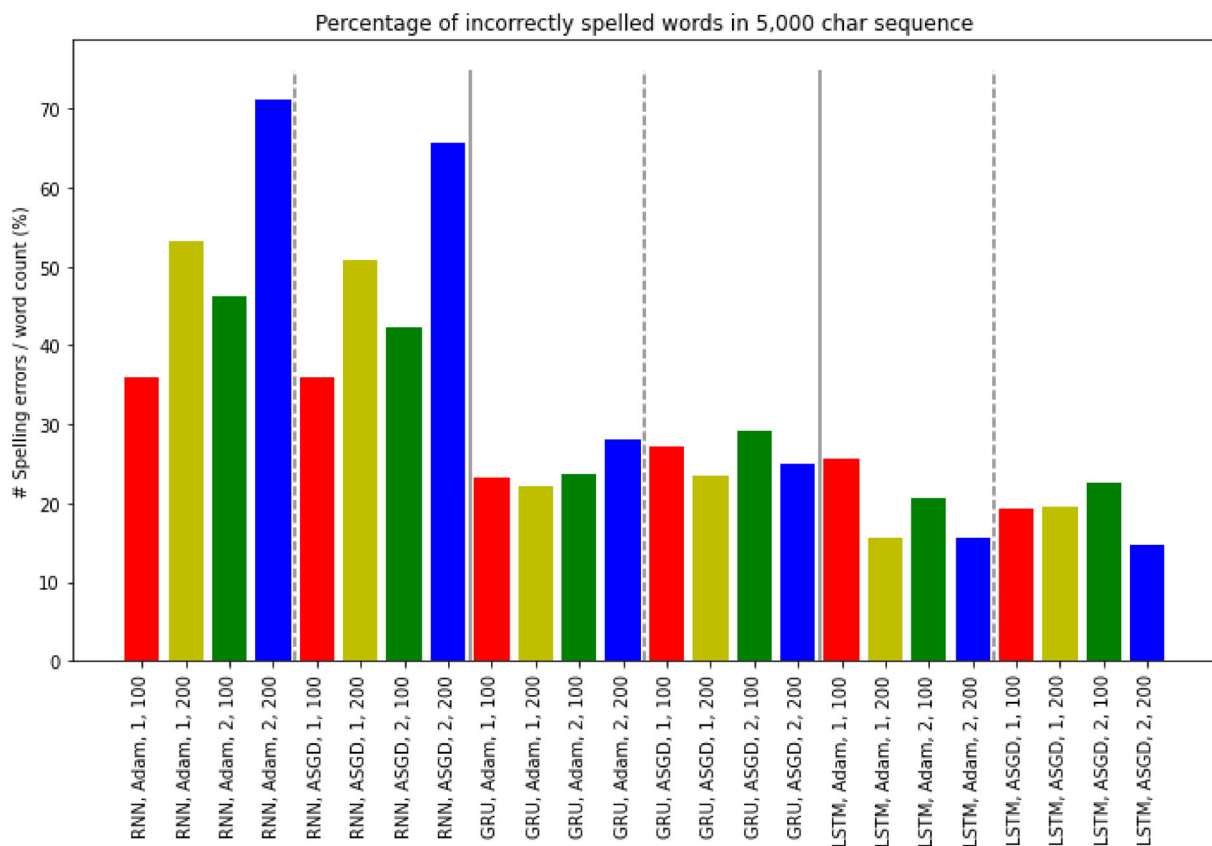
print('Shortest training time: LSTM, Adam, 1, 100. ({} min)'.format(df['training
```



Shortest training time: LSTM, Adam, 1, 100. (26.35854196 min)

```
In [12]: errors = df['% spelling errors'].to_numpy()*100
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(times)),errors,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], 0, 75, 'gray')
plt.vlines([3.5,11.5,19.5], 0, 75, 'gray','dashed')
plt.ylabel('# Spelling errors / word count (%)')
plt.title('Percentage of incorrectly spelled words in 5,000 char sequence')
plt.tight_layout()
plt.savefig('Images/spelling_percentage.png')
plt.show()

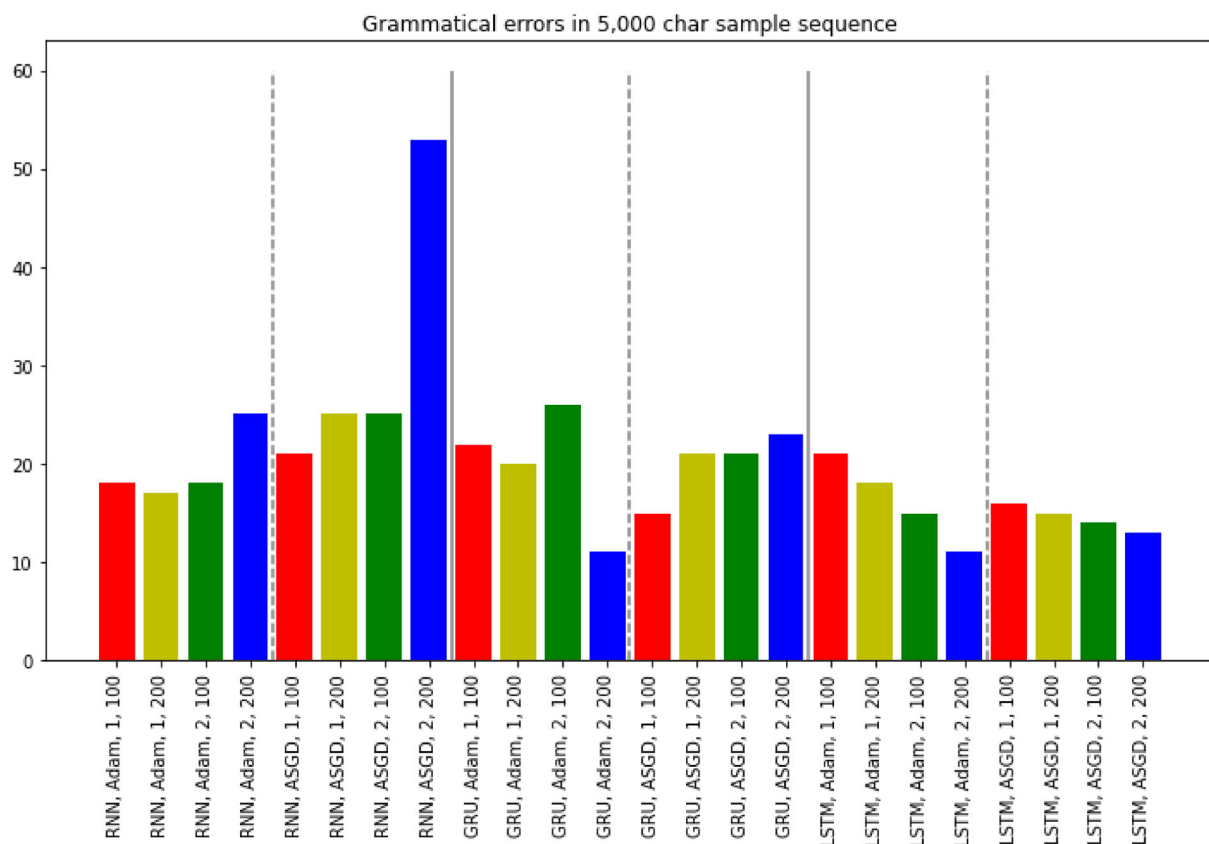
print('Smallest percentage of incorrectly spelled words: LSTM, ASGD, 2, 200. ({}%
```



Smallest percentage of incorrectly spelled words: LSTM, ASGD, 2, 200. (14.7161066%)


```
In [13]: gerrors = df['grammer errors'].to_numpy()
plt.figure(figsize=(10,7))
plt.bar(np.arange(len(times)),gerrors,color=['r','y','g','b'],tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], 0, 60, 'gray')
plt.vlines([3.5,11.5,19.5], 0, 60, 'gray','dashed')
plt.title('Grammatical errors in 5,000 char sample sequence')
plt.tight_layout()
plt.savefig('Images/grammar_errors.png')
plt.show()

print('Fewest grammatical errors: GRU, Adam, 2, 200. ({} )'.format(df['grammer err
```



Fewest grammatical errors: GRU, Adam, 2, 200. (11)

```
In [14]: data = (df['Sample word count'].to_numpy() - n1)**-1
plt.figure(figsize=(10,7))
colors=['g' if i>=0 else 'r' for i in data]
plt.bar(np.arange(len(times)),data,color=colors,tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], data.min(), data.max(), 'gray')
plt.vlines([3.5,11.5,19.5], data.min(), data.max(), 'gray', 'dashed')
plt.ylabel('1 / (Sample word count - expected word count)')
plt.title('Inverse difference between expected and actual word count (per 5,000 c
plt.tight_layout()
plt.savefig('Images/word_count_comparison.png')
plt.show()
```



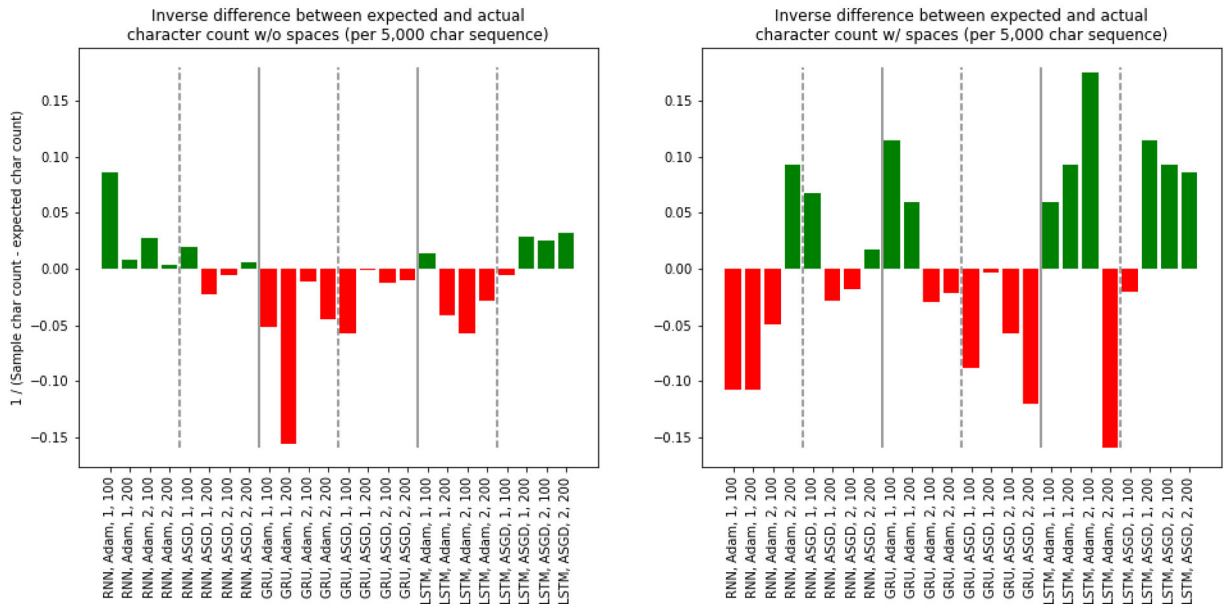
```
In [15]: plt.figure(figsize=(16,6))
plt.subplot(1,2,1)
data = (df['Sample char count w/o spaces'].to_numpy() - n3)**-1

colors=['g' if i>=0 else 'r' for i in data]
plt.bar(np.arange(len(times)),data,color=colors,tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], -0.16, 0.18, 'gray')
plt.vlines([3.5,11.5,19.5], -0.16, 0.18, 'gray','dashed')
plt.ylabel('1 / (Sample char count - expected char count)')
plt.title('Inverse difference between expected and actual\ncharacter count w/o spaces')

plt.subplot(1,2,2)
data = (df['Sample char count w/ spaces'].to_numpy() - n2)**-1

colors=['g' if i>=0 else 'r' for i in data]
plt.bar(np.arange(len(times)),data,color=colors,tick_label=labels)
plt.xticks(rotation='vertical')
plt.vlines([7.5,15.5], -0.16, 0.18, 'gray')
plt.vlines([3.5,11.5,19.5], -0.16, 0.18, 'gray','dashed')
plt.title('Inverse difference between expected and actual\ncharacter count w/ spaces')

plt.savefig('Images/char_count_comparison.png')
plt.show()
```



In [ ]:

