

```

1 ##### DESCRIPTION #####
2 # Use to train a model on the specified dataset and hyperparameters.
3
4 import string
5 import random
6 import torch
7 import torch.nn as nn
8 import matplotlib.pyplot as plt
9 import numpy as np
10 import time
11 import sys
12 import os
13 from tqdm import tqdm
14
15 FILE_NAME = 'complete_sherlock_holmes.txt'
16 DATASET = 'Complete Sherlock Holmes'
17
18 ##### HYPERPARAMETERS #####
19 CELL_TYPE = 'RNN' # DEFAULTS TO 'RNN'. Options: ['RNN', 'LSTM', 'GRU', 'RELU']
20 OPTIM_TYPE = 'Adam' # DEFAULTS TO 'Adam'. Options: ['Adam', 'ASGD', 'Adagrad', 'RMSprop']
21 HIDDEN_LAYERS = 1 # DEFAULT: 1
22 HIDDEN_SIZE = 100 # DEFAULT: 100
23
24 LEARNING_RATE = 0.005 # 0.005 for Adam, 0.05 for other optimizers
25 INPUT_SEQUENCE = 128 # DEFAULT: 128
26
27 TRAINING_ITERATIONS = 20000 # DEFAULT: 20000
28 INITIAL_SEQUENCE = '\n' # To use for eval
29
30 print('Running on', CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE)
31
32 all_chars = string.printable
33 n_chars = len(all_chars)
34 file = open('./'+FILE_NAME).read()
35 file_len = len(file)
36
37 def get_random_seq():
38     seq_len = INPUT_SEQUENCE # The length of an input sequence.
39     start_index = random.randint(0, file_len - seq_len)
40     end_index = start_index + seq_len + 1
41     return file[start_index:end_index]
42 def seq_to_onehot(seq):
43     tensor = torch.zeros(len(seq), 1, n_chars)
44     for t, char in enumerate(seq):
45         index = all_chars.index(char)
46         tensor[t][0][index] = 1
47     return tensor
48 def seq_to_index(seq):
49     tensor = torch.zeros(len(seq), 1)
50     for t, char in enumerate(seq):
51         tensor[t] = all_chars.index(char)
52     return tensor
53 def get_input_and_target():
54     seq = get_random_seq()
55     input = seq_to_onehot(seq[:-1]) # Input is represented in one-hot.
56     target = seq_to_index(seq[1:]).long() # Target is represented in index.

```

```

57     return input, target
58
59 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
60 print('Device =', device)
61
62 class Net(nn.Module):
63     def __init__(self):
64         # Initialization.
65         super(Net, self).__init__()
66         self.input_size = n_chars # Input size: Number of unique chars.
67         self.hidden_size = HIDDEN_SIZE
68         self.output_size = n_chars # Output size: Number of unique chars.
69
70         # Ensures the size of the hidden layer stack does not exceed 3
71         self.layers = HIDDEN_LAYERS
72         if HIDDEN_LAYERS>3:
73             self.layers=3
74
75         # Create a rnn cell for the stack
76         def create_cell(size_in, size_out):
77             if CELL_TYPE=='LSTM':
78                 return nn.LSTMCell(size_in, size_out)
79             elif CELL_TYPE=='GRU':
80                 return nn.GRUCell(size_in, size_out)
81             elif CELL_TYPE=='RELU': # Not used in testing report
82                 return nn.RNNCell(size_in, size_out, nonlinearity='relu')
83             else:
84                 return nn.RNNCell(size_in, size_out)
85
86         self.rnn = create_cell(self.input_size, self.hidden_size)
87
88         if HIDDEN_LAYERS>=2:
89             self.rnn2 = create_cell(self.hidden_size, self.hidden_size)
90         if HIDDEN_LAYERS>=3:
91             self.rnn3 = create_cell(self.hidden_size, self.hidden_size)
92
93         self.fc = nn.Linear(self.hidden_size, self.output_size)
94
95     def forward(self, input, hidden, cell, hidden2=False, cell2=False, hidden3=False,
96 cell3=False):
97         # Forward function.
98         # takes in the 'input' and 'hidden' tensors,
99         # can also take in 'cell state' tensor if cell type is 'LSTM',
100        # takes additional hidden and cell state tensors for each layer
101        if CELL_TYPE=='LSTM':
102            hidden, cell = self.rnn(input, (hidden,cell))
103            if self.layers>=2:
104                hidden2, cell2 = self.rnn2(hidden, (hidden2,cell2))
105            if self.layers>=3:
106                hidden3, cell3 = self.rnn3(hidden2, (hidden3,cell3))
107        else:
108            hidden = self.rnn(input, hidden)
109            if self.layers>=2:
110                hidden2 = self.rnn2(hidden, hidden2)
111            if self.layers>=3:
112                hidden3 = self.rnn3(hidden2, hidden3)

```

```

113     # Linear transformation (fully connected layer) to the output
114     if self.layers==3:
115         output = self.fc(hidden3)
116         return output, hidden, cell, hidden2, cell2, hidden3, cell3
117     elif self.layers==2:
118         output = self.fc(hidden2)
119         return output, hidden, cell, hidden2, cell2
120     else:
121         output = self.fc(hidden)
122         return output, hidden, cell
123 def init_hidden(self):
124     # Initial hidden state.
125     return torch.zeros(1, self.hidden_size).to(device)
126 def init_cell(self):
127     # Initial cell state.
128     return torch.zeros(1, self.hidden_size).to(device)
129
130 net = Net()
131 net.to(device)
132
133 # Training step function
134 def train_step(net, opt, input, target):
135     seq_len = input.shape[0]
136     hidden = net.init_hidden() # Initial hidden state
137     cell = net.init_cell() # Initial cell state
138     if HIDDEN_LAYERS >=2:
139         hidden2 = net.init_hidden()
140         cell2 = net.init_cell()
141     if HIDDEN_LAYERS >=3:
142         hidden3 = net.init_hidden()
143         cell3 = net.init_cell()
144
145     net.zero_grad()
146     loss = 0 # Initial loss.
147
148     for t in range(seq_len): # For each one in the input sequence
149         if HIDDEN_LAYERS==3:
150             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(input[t], hidden,
cell, hidden2, cell2, hidden3, cell3)
151         elif HIDDEN_LAYERS==2:
152             output, hidden, cell, hidden2, cell2 = net(input[t], hidden, cell, hidden2,
cell2)
153         else:
154             output, hidden, cell = net(input[t], hidden, cell)
155         loss += loss_func(output, target[t])
156
157     loss.backward() # Backward.
158     opt.step() # Update the weights.
159
160     return loss / seq_len
161
162 # Evaluation step function
163 def eval_step(net, init_seq=INITIAL_SEQUENCE, predicted_len=100):
164     # Initialize the hidden state, input and the predicted sequence
165     hidden = net.init_hidden()
166     cell = net.init_cell()
167     if HIDDEN_LAYERS >=2:

```

```

168         hidden2 = net.init_hidden()
169         cell2 = net.init_cell()
170     if HIDDEN_LAYERS >= 3:
171         hidden3 = net.init_hidden()
172         cell3 = net.init_cell()
173     init_input = seq_to_onehot(init_seq).to(device)
174     predicted_seq = init_seq
175
176     # Use initial string to "build up" hidden state.
177     for t in range(len(init_seq) - 1):
178         if HIDDEN_LAYERS==3:
179             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(init_input[t], hidden,
cell, hidden2, cell2, hidden3, cell3)
180         elif HIDDEN_LAYERS==2:
181             output, hidden, cell, hidden2, cell2 = net(init_input[t], hidden, cell, hidden2,
cell2)
182         else:
183             output, hidden, cell = net(init_input[t], hidden, cell)
184         # Set current input as the last character of the initial string.
185         input = init_input[-1]
186
187     # Predict more characters after the initial string.
188     for t in range(predicted_len):
189         # Get the current output and hidden state.
190         if HIDDEN_LAYERS==3:
191             output, hidden, cell, hidden2, cell2, hidden3, cell3 = net(input, hidden, cell,
hidden2, cell2, hidden3, cell3)
192         elif HIDDEN_LAYERS==2:
193             output, hidden, cell, hidden2, cell2 = net(input, hidden, cell, hidden2, cell2)
194         else:
195             output, hidden, cell = net(input, hidden, cell)
196
197     # Sample from the output as a multinomial distribution.
198     try:
199         predicted_index = torch.multinomial(output.view(-1).exp(), 1)[0]
200     except: # Added post to resolve errors with tensors containing 'inf'/'nan' values
201         predicted_index = torch.multinomial(output.view(-1).exp().clamp(0.0, 3.4e38), 1)[0]
202
203     # Add predicted character to the sequence and use it as next input.
204     predicted_char = all_chars[predicted_index]
205     predicted_seq += predicted_char
206     # Use the predicted character to generate the input of next round.
207     input = seq_to_onehot(predicted_char)[0].to(device)
208
209     return predicted_seq
210
211 ##### MAIN ALGORITHM #####
212
213 iters = TRAINING_ITERATIONS # Number of training iterations.
214
215 # The loss variables.
216 all_losses = []
217 # Initialize the optimizer and the loss function.
218 if(OPTIM_TYPE=='ASGD'):
219     opt=torch.optim.ASGD(net.parameters(), lr=LEARNING_RATE)
220 if(OPTIM_TYPE=='Adagrad'): # Not used in testing results

```

```

221     opt=torch.optim.Adagrad(net.parameters(), lr=LEARNING_RATE)
222     if(OPTIM_TYPE=='RMSprop'): # Not used in testing results
223         opt=torch.optim.RMSprop(net.parameters(), lr=LEARNING_RATE)
224     else:
225         opt = torch.optim.Adam(net.parameters(), lr=LEARNING_RATE)
226     loss_func = nn.CrossEntropyLoss()
227
228     # Training procedure.
229     start_time = time.time()
230     for i in tqdm(range(iters)):
231         input, target = get_input_and_target() # Fetch input and target.
232         input, target = input.to(device), target.to(device) # Move to GPU memory.
233         loss = train_step(net, opt, input, target) # Calculate the loss.
234         all_losses.append(loss)
235
236     end_time = time.time()
237     total_time = end_time - start_time
238
239     # Calculates summary of losses
240     i_half = int(len(all_losses)*0.5)
241     i_quart = int(len(all_losses)*0.75)
242     loss_avg = np.sum(np.array(all_losses))/len(all_losses)
243     loss_avg_half = np.sum(np.array(all_losses[i_half:]))/len(all_losses[i_half:])
244     loss_avg_quart = np.sum(np.array(all_losses[i_quart:]))/len(all_losses[i_quart:])
245     loss_list=[elem.item() for elem in [loss_avg,loss_avg_half,loss_avg_quart]]
246     rolling_losses=[]
247     losses_copy = [i.item() for i in all_losses]
248     for i in range(len(losses_copy)):
249         temp=losses_copy[np.max((i-100, 0)):i+1]
250         rolling_losses.append(np.sum(temp)/len(temp))
251
252     plt.xlabel('iters')
253     plt.ylabel('loss')
254     plt.hlines(loss_list,0,len(losses_copy)-1,['red','orange','green'],'dashed')
255     plt.plot(rolling_losses)
256     plt.ylim(0,5)
257
258     print('Avg loss: {}'.format(loss_avg))
259     print('Avg loss last half: {}'.format(loss_avg_half))
260     print('Avg loss last quarter: {}'.format(loss_avg_quart))
261     print()
262     print('Training time: {} sec'.format(total_time))
263     print('{} min | {:.3f} hr'.format(total_time/60, total_time/3600))
264
265     # Creates a folder path to save training results to
266     PATH = 'Results'
267     for folder in [DATASET, CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE]:
268         folder=str(folder)
269         if not os.path.isdir(PATH+'/'+folder):
270             os.mkdir(PATH+'/'+folder)
271         PATH+='/' + folder
272     PATH+='/'
273
274     print()
275     print('Results saved to: {}'.format(PATH))
276     model_path = PATH+'model.pt' # Saves model parameters
277

```

```
278 torch.save(net.state_dict(), model_path)
279
280 # Sequence of all 20000 loss values
281 file = open(PATH+'all_losses.txt', 'w')
282 file.write(' '.join([str(elem) for elem in losses_copy]))
283 file.close()
284
285 # A 5000 char sample generated after training
286 file = open(PATH+'sample.txt', 'w')
287 file.write(eval_step(net, predicted_len=5000))
288 file.close()
289
290 # Information on training
291 file = open(PATH+'info.txt', 'w')
292 file.write('Iterations: {}\n\n'.format(TRAINING_ITERATIONS))
293 file.write('Dataset: {}\nInput Size: {}\nLearning Rate: {}\n\n'.format(DATASET,
    INPUT_SEQUENCE, LEARNING_RATE))
294 file.write('Cell Type: {}\nOptimizer: {}\nHidden Layers: {}\nHidden Size:
    {}\n\n'.format(CELL_TYPE, OPTIM_TYPE, HIDDEN_LAYERS, HIDDEN_SIZE))
```